

ETL Project Documentation — Movies Data ETL using Python & Docker

Project Overview

This project involves setting up an ETL (Extract, Transform, Load) pipeline for movie data using Docker, Google Cloud Pub/Sub emulator, and Python scripts. The primary goal is to process and analyze movie rating data, leveraging containerization and Pub/Sub for efficient data handling and processing.

Detailed Explanation

1. Environment Setup

setup_env.bat

```
1. @echo off
2. py -m venv venv
3. call venv\Scripts\activate
4. pip install numpy pandas google-cloud-pubsub docker
5. echo Environment setup completed.
```

- py -m venv venv: Creates a virtual environment named venv to manage dependencies.
- call venv\Scripts\activate: Activates the virtual environment.
- pip install numpy pandas google-cloud-pubsub docker: Installs the required Python libraries:
 - o numpy and pandas for data manipulation.
 - o google-cloud-pubsub for interacting with Google Pub/Sub.
 - o docker for Docker client interactions.
- echo Environment setup completed.: Provides feedback that the setup is complete.

Why use a virtual environment?

- To isolate project dependencies from the global Python environment and avoid conflicts.

2. Data Download

download_data.bat

```
1. @echo off
2. REM Create the data directory if it doesn't exist
3. if not exist "..\data" mkdir ..\data
4.
5. echo Downloading MovieLens dataset...
6. curl -o ..\data\ml-100k.zip https://files.grouplens.org/datasets/movielens/ml-100k.zip
7. echo Unzipping dataset...
8. powershell -command "Expand-Archive -Path '..\data\ml-100k.zip' -DestinationPath '..\data\'
9. echo Data download and extraction complete.
```

- If not exist "..\data" mkdir ..\data: Checks if the data directory exists; if not, it creates it.
- curl -o ..\data\ml-100k.zip https://files.grouplens.org/datasets/movielens/ml-100k.zip: Downloads the MovieLens dataset zip file.
- powershell -command "Expand-Archive -Path '..\data\ml-100k.zip' -DestinationPath '..\data\'": Unzips the dataset using PowerShell.

Why download and unzip the dataset?

- To prepare the raw data for further processing.

3. Dockerfile

```
1. Use the official Python image from the Docker Hub
2. FROM python:3.12.5-slim
3.
4. # Set the working directory in the container
5. WORKDIR /app
6. # Copy the current directory contents into the container at /app
7. COPY . /app
8.
9. # Install the required libraries
10. RUN pip install --no-cache-dir numpy pandas google-cloud-pubsub scipy matplotlib seaborn subprocess
11.
12. # Command to start the Pub/Sub emulator (if needed)
13. CMD ["gcloud", "beta", "emulators", "pubsub", "start", "--host-port=0.0.0.0:8085"]
```

- FROM python: 3.12.5-slim: Uses a lightweight Python image to build the Docker container.
- WORKDIR /app: Sets the working directory inside the container.
- COPY . /app: Copies all files from the host to the container.
- RUN pip install --no-cache-dir ...: Installs required libraries.
- CMD ["gcloud", "beta", "emulators", "pubsub", "start", "--host-port=0.0.0.0:8085"]: Starts the Pub/Sub emulator within the container.

Why use Docker?

- Ensures a consistent environment for running and testing your application.

4. Start Pub/Sub Emulator

start_emulator.py

```
1. import os
2. # Start the Pub/Sub emulator in a Docker container
3. os.system("docker run -d --name pubsub-emulator -p 8085:8085 google/cloud-sdk gcloud beta emulators pubsub
start --host-port=0.0.0.0:8085")
4. print("Pub/Sub emulator started.")
```

- os.system(...): Executes a command to start the Pub/Sub emulator in a Docker container.

Why use an emulator?

- Allows for local development and testing of Pub/Sub functionality without incurring cloud costs.

5. Build Docker Image

```
1. docker build -t movies-etl-image .
```

- Builds a Docker image named movies-etl-image based on the Dockerfile.

Why build an image?

- To package the application and its dependencies into a single unit for easy deployment.

6. Run Docker Container

```
1. docker run -d --name movies-etl-container -p 8085:8085 movies-etl-image
```

- Runs a Docker container named movies-etl-container based on the built image.

Why run a container?

- To execute the application in a controlled environment.

7. Create Topic and Subscription

create_topic_subscription.py

```
1. import os
2. from google.cloud import pubsub_v1
3.
4. # Set up the environment variable to point to the local Pub/Sub emulator
5. os.environ["PUBSUB_EMULATOR_HOST"] = "localhost:8085"
6.
7. # Set the project ID and other parameters
8. project_id = "test-project"
9. topic_id = "movies-topic"
10. subscription_id = "movies-subscription"
11.
12. # Set up Pub/Sub client
13. publisher = pubsub_v1.PublisherClient()
14. subscriber = pubsub_v1.SubscriberClient()
15.
16. # Create a topic
17. project_path = f"projects/{project_id}"
18. topic_path = publisher.topic_path(project_id, topic_id)
19. publisher.create_topic(name=topic_path)
20. print(f"Topic '{topic_id}' created.")
21.
22. # Create a subscription
23. subscription_path = subscriber.subscription_path(project_id, subscription_id)
24. subscriber.create_subscription(name=subscription_path, topic=topic_path)
25. print(f"Subscription '{subscription_id}' created.")
```

- os.environ["PUBSUB_EMULATOR_HOST"] = "localhost:8085": Sets the environment variable to use the local Pub/Sub emulator.
- publisher.create_topic(...): Creates a Pub/Sub topic.
- subscriber.create_subscription(...): Creates a subscription for the topic.

Why create topics and subscriptions?

- To organize and manage data flows within the Pub/Sub system.

8. Extract CSV Files

process_data.py

```
1. import pandas as pd
2.
3. # Define the path to the data files
4. ratings_file = 'D:/DEPI Data Engineer/0_projects/ETL_MOVIES/venv/data/ml-100k/u.data'
5. movies_file = 'D:/DEPI Data Engineer/0_projects/ETL_MOVIES/venv/data/ml-100k/u.item'
6.
7. # Create a DataFrame for Movie Ratings
8. ratings = pd.read_csv(ratings_file, delimiter='\t', header=None, names=['user_id', 'item_id', 'rating', 'timestamp'])
9.
10. # Create a DataFrame for Movie Names
11. with open(movies_file, 'r', encoding="ISO-8859-1") as read_file:
12.     counter = 0
13.     movies_df = pd.DataFrame(columns=['item_id', 'movie_name', 'release_timestamp'])
14.
15.     # Iterate through the lines in the file
16.     for line in read_file:
17.         fields = line.split('|')
18.         item_id, movie_name, release_timestamp = fields[0], fields[1], fields[2]
19.         movie_name = movie_name[0:len(movie_name) - len(' (1234)')]
20.
```

```

21.         line_data = [int(item_id), str(movie_name), release_timestamp]
22.
23.         temp_df = pd.DataFrame(data=[line_data], columns=['item_id', 'movie_name', 'release_timestamp'])
24.         movies_df = pd.concat([temp_df, movies_df], ignore_index=True)
25.
26.         counter += 1
27.
28.         # Sort Values by item id
29.         movies_df.sort_values(by='item_id', ascending=True, inplace=True)
30.
31. # Close file
32. read_file.close()
33.
34. # Export to CSV
35. ratings.to_csv('D:/DEPI Data Engineer/0_projects/ETL_MOVIES/venv/data/ratings.csv', index=False)
36. movies_df.to_csv('D:/DEPI Data Engineer/0_projects/ETL_MOVIES/venv/data/movies.csv', index=False)

```

- pd.read_csv(...): Reads the movie ratings and movie names data into DataFrames.
- with open(...) as read_file: Reads movie names file and processes each line.
- pd.concat(...): Aggregates movie data into a single DataFrame.
- ratings.to_csv(...) and movies_df.to_csv(...): Saves the DataFrames as CSV files.

Why extract and save data?

- To prepare the raw data for preprocessing and analysis.

9. Preprocess Data

preprocessing_data.py

```

1. import pandas as pd
2. import os
3. import numpy as np
4. from scipy import stats
5. import seaborn as sns
6. import matplotlib.pyplot as plt
7.
8. # Define paths for the existing files
9. ratings_csv_path = r'D:\DEPI Data Engineer\0_projects\ETL_MOVIES\venv\data\ratings.csv'
10. movies_csv_path = r'D:\DEPI Data Engineer\0_projects\ETL_MOVIES\venv\data\movies.csv'
11. merged_csv_path = r'D:\DEPI Data Engineer\0_projects\ETL_MOVIES\venv\data\full_data.csv'
12.
13. # Load the datasets into DataFrames
14. ratings_df = pd.read_csv(ratings_csv_path)
15. movies_df = pd.read_csv(movies_csv_path)
16.
17. full_data = pd.merge(ratings_df, movies_df, on='item_id')
18.
19. # Save the merged DataFrame to a new CSV file
20. full_data.head(10)
21. full_data = full_data.sort_values(by='item_id', ascending=True)
22. full_data
23.
24. # data preprocessing
25. ### 1st step: check data:
26. full_data.info()
27. missing_percentage = full_data['release_timestamp'].isnull().mean() * 100
28. print(f"Missing percentage: {missing_percentage:.2f}%")
29. mode_timestamp = full_data['release_timestamp'].mode()[0]
30. full_data['release_timestamp'].fillna(mode_timestamp, inplace=True)
31. print(full_data['release_timestamp'].isnull().sum())
32.
33. # Check outliers
34. def detect_outliers(data, column):
35.     z_scores = np.abs(stats.zscore(data[column]))
36.     return np.where(z_scores > 3)
37.
38. sns.boxplot(x=full_data['rating'])
39.
40. full_data['timestamp'] = pd.to_datetime(full_data['timestamp'], unit='s')
41. full_data['release_timestamp'] = pd.to_datetime(full_data['release_timestamp'], format='%d-%b-%Y')

```

```

42. full_data['datestamp'] = full_data['timestamp'].dt.date
43. full_data['time'] = full_data['timestamp'].dt.time
44. full_data = full_data.drop(columns=['timestamp'])
45.
46. invalid_rows = full_data[full_data['datestamp'] < full_data['release_timestamp']]
47. unique_invalid_movies = invalid_rows['movie_name'].unique()
48. plt.figure(figsize=(15, 5))
49. sns.countplot(x=invalid_rows['movie_name'])
50. plt.xticks(rotation=90, fontsize=8)
51. plt.show()
52.
53. len(full_data[full_data.duplicated()])
54. full_data.to_csv(merged_csv_path, index=False)

```

- `pd.merge(...)`: Merges ratings and movie names DataFrames.
- `full_data.info()`: Displays info about the DataFrame to check for missing values and data types.
- `full_data['release_timestamp'].fillna(mode_timestamp, inplace=True)`: Fills missing values in the `release_timestamp` column.
- `detect_outliers(...)`: Function to detect outliers in data.
- `sns.boxplot(...)`: Creates a boxplot to visualize outliers.
- `pd.to_datetime(...)`: Converts timestamps to datetime objects.
- `full_data.drop(columns=['timestamp'])`: Drops the original timestamp column.
- `full_data[full_data['datestamp'] < full_data['release_timestamp']]`: Identifies invalid rows where the release date is after the rating date.
- `plt.figure(...)`: Creates a bar plot of invalid movies.
- `full_data.to_csv(...)`: Saves the cleaned DataFrame to a CSV file.

Why preprocess data?

- To clean and prepare the data for analysis, ensuring consistency and quality.

10. Create Folder in Docker Container

```
1. mkdir -p /data
```

- Creates a `/data` directory in the running Docker container.
- Run this command using terminal of your container inside docker

Why create a directory in Docker?

- To store and access data files within the container.

11. Publish Data to Docker Container

`publish_data.py`

```

1. import subprocess
2.
3. # Paths for the files and Docker container
4. file_path = r'D:\DEPI Data Engineer\0_projects\ETL_MOVIES\env\data\full_data.csv'
5. container_id = '54fea1529a3c1360e7a9be6f5711b34b7cf6cb05955e99cba5aa59ffd1e6c95b'
6. container_path = '/data/full_data.csv'
7.
8. def copy_to_container(file_path, container_id, container_path):
9.     try:
10.         result = subprocess.run(['docker', 'cp', file_path, f'{container_id}:{container_path}'],
11.                                 stdout=subprocess.PIPE, stderr=subprocess.PIPE, text=True)
12.
13.         if result.returncode == 0:
14.             print(f"Successfully copied {file_path} to {container_id}:{container_path}")
15.         else:
16.             print(f"Error copying file: {result.stderr}")
17.             return False
18.         return True

```

```

19.     except Exception as e:
20.         print(f"Exception occurred: {e}")
21.         return False
22.
23. # Copy the CSV file to the Docker container
24. copy_to_container(file_path, container_id, container_path)

```

- subprocess.run(...): Executes a Docker command to copy the CSV file into the container.
- docker cp: Copies files between the host and Docker containers.

Why copy data to Docker?

- To make the processed data available inside the Docker container for further operations.

Tools and Alternatives

1. Python & Libraries:

- **pandas**: For data manipulation.
- **numpy**: For numerical operations.
- **google-cloud-pubsub**: For interacting with Pub/Sub.

2. Docker:

- For Containerization: Ensures environment consistency.

3. Pub/Sub Emulator:

- For Local Testing: Avoids cloud costs.

Sending Data: Messages vs. CSV

	Messages	CSV Files
Pros	Efficient for real-time processing; no need to handle large files.	Simpler to handle; compatible with many systems.
Cons	Requires parsing and structuring messages.	Simpler to handle; compatible with many systems.

Recommendation:

- For your current setup, sending data as CSV files might be simpler and more compatible with existing tools. Messages could be considered for real-time or streaming scenarios.
- prepare the data for analysis, ensuring consistency and quality.

References:

- [Data Engineering Project — Movies Data ETL using Python & GCP](#)
- [Running GCP Pub/Sub Emulator on a Local Docker Environment](#)

About The Author:

- **Name:** Nada Hamdy Fatehy
- **Email:** nadahamdy2172002@gmail.com
- **Phone:** +201207788890
- **Links:** [LinkedIn](#) | [GitHub](#)